

Human-level control through deep reinforcement learning



Jake Hoare, UNSW
May 2016

Objective

- Create a single agent that can play a range of different Atari games better than a human
 - Inputs are screen pixels and score
 - Actions are valid moves in the game
 - General AI
-
-

Challenges

- Natural solution – reinforcement learning
 - High dimensionality of input – 210 x 160 screen with 128 colours at 60 Hz
 - No domain-specific information or heuristics
 - Delayed rewards
 - Instability of reinforcement learning with non-linear approximator
-
-

Q-learning

- Learn the policy (what action to take in a given state) that maximises the sum of future discounted rewards

$$Q^*(s,a) = \max_{\pi} \mathbb{E} [r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = a, \pi]$$

$$Q^*(s,a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q^*(s',a') | s,a \right]$$

Dimensionality


- State becomes sequence of 4 frames
- $10^{283,000}$ states
- Need to generalise to reduce dimensions
- Instead of learning discrete $Q(s,a)$, approximate with convolutional neural network

$$Q(s,a; \theta) \approx Q^*(s,a)$$

Initial data compression

- Colour 210 x 160 to 256-grayscale 84 x 84
- Max(frame, previous frame)

Network architecture



- Learn a separate Q for each action – find the best action by choosing the largest

Experience replay

- Instability of network learning caused by temporal correlations in sequence of observations
- Network trained on random minibatches of previous observations
- Efficient reuse of experiences

Freeze target

- Training example is a tuple (s,a,r,s')
- Error function uses target value also based on network output
- Poor convergence if output and target move together so adjust target network only periodically

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

rmsProp

- Problem of same gradient descent learning rate for all weights
 - rprop – for full batch learning, increase the step size multiplicatively if gradients agree
 - rmsProp – for minibatch learning, normalise update by maintaining a decaying mean square average of each gradient
-
-

Other tricks

- Reward clipping – normalise to +/-1, loses ability to differentiate between big and small gains
 - Skipping frames – fast-forward by selecting actions every 4th frame
 - Exploration vs exploitation – epsilon greedy policy to explore initially but then take best actions
-
-

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$



 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

Results



Strengths and weaknesses

- Space Invaders
https://youtu.be/Dds_yDJFhvI
 - Breakout
<https://youtu.be/cjpElotvwFY>
 - Montezuma
<https://www.youtube.com/watch?v=Klxxg9JM5tY>
-
-

Double Q learning

- Improve upon target freezing by unpacking the target expression
- Network used twice – once to determine the best next action and once to find the Q value of the next state

$$Y_t^{\text{DoubleQ}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t); \theta'_t)$$

Prioritised experience replay

- Improve upon experience replay by replaying large error examples more often
- Can overfit and may not replay if error was initially low – add randomness